

causalscale: A Unified Causal Discovery Engine with Pre-trained Backbones

Shuaidong Gao

gaoshuaidong@stu.cqifs.edu.cn

Chongqing Institute of Foreign Studies

Qijiang, China

Abstract

We present *causalscale* (v3.1.0), an open-source Python package unifying causal discovery via six engines—verified NOTEARS, LowRankGNN, MultiScale, ensemble consensus, bootstrap stability, and ASCEND two-tier. On ER DAGs ($d=30-100$), *causalscale* achieves $F_1=0.586-0.462$ vs. NOTEARS $F_1=0.581-0.185$, with the advantage widening from +1% at $d=30$ to +150% at $d=100$. A 33-cancer scan recovers ARID1A–MTOR directionality. On STRING-anchored DepMap ($d=200$), ASCEND two-tier achieves 93.3% STRING/TRRUST validation precision (14/15 edges). Features: four-mode `validate()` API, ensemble consensus voting, bootstrap stability. Code: <https://github.com/sgao-academics/causalscale> (MIT).

CCS Concepts

• **Computing methodologies** → **Causal reasoning and diagnostics**; *Machine learning approaches*; • **Applied computing** → *Computational biology*.

Keywords

causal discovery, open-source software, benchmarking, NOTEARS, cancer genomics

ACM Reference Format:

Shuaidong Gao. 2026. *causalscale: A Unified Causal Discovery Engine with Pre-trained Backbones*. In . ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 Introduction

Causal discovery from observational data is a fundamental problem spanning machine learning [13], statistics [9], and biomedical science [6]. NOTEARS [13] reformulated the combinatorial DAG constraint as the smooth function $h(W) = \text{tr}(\exp(W \circ W)) - d = 0$, enabling gradient-based structure learning via augmented Lagrangian optimization. This spawned DAGMA [1] (log-determinant acyclicity), Golem [8] (likelihood-based objectives), and nonlinear extensions. Yet all share an $O(d^3)$ matrix-exponential bottleneck, limiting deployment to $d \leq 150$ [5].

Domain scientists face a stark practical choice: select a handful of candidate genes or sacrifice causal directionality for scale. Genomics researchers cannot run DAG-constrained discovery at genome-wide resolution ($d \approx 20,000$). Neuroscientists studying fMRI connectomics ($d \approx 10^4$) lack DAG-compatible tools. Financial analysts hit the 150-asset ceiling. This impasse is compounded by software fragmentation: each existing library (DoWhy, gCastle,

causal-learn, GENIE3) has separate installation, separate APIs, and no path to genome-scale inference.

LowRankGNN [4] broke the $O(d^3)$ barrier via $W = UV^\top$ factorization, achieving 93.3% STRING/TRRUST validation precision via ASCEND two-tier discovery, but existed only as a research prototype. *causalscale* (v3.1.0) addresses this gap as the first production-grade causal discovery package that: (1) scales from $d = 30$ to 10^8 on a single consumer GPU; (2) unifies six engines under automatic selection; (3) ships with pre-trained backbones for zero-shot transfer; (4) provides four-mode `validate()` API; (5) supports ensemble consensus and bootstrap stability; and (6) includes ASCEND two-tier biological discovery. It is accompanied by synthetic benchmarks ($d = 30-100$), hyperparameter ablations, STRING/TRRUST-validated genome-scale validation, and a 33-cancer application.

2 Related Work

2.1 Differentiable Causal Discovery

NOTEARS [13] introduced the continuous DAG constraint, transforming combinatorial search into constrained optimization. Golem [8] replaced the least-squares objective with a likelihood-based formulation, improving edge recovery at small d but retaining $O(d^3)$ complexity. DAGMA [1] introduced log-determinant acyclicity $\log \det(sI - W \circ W) - d \log s$, theoretically tighter than NOTEARS’ $h(W)$ but exhibiting severe convergence pathology—producing zero edges at all tested dimensions in our benchmarks under default hyperparameters. The ρ -Diagnostic framework [5] systematically characterized death-line failure modes at $d \approx 150$ where the L1 penalty dominates the reconstruction gradient, causing $W \rightarrow 0$. LowRankGNN [4] circumvented this bottleneck through low-rank factorization, reducing complexity to $O(dr^2)$ with a specialized optimization procedure, but lacked software packaging.

2.2 Causal Discovery Software Landscape

DoWhy [10] targets causal inference (effect estimation and refutation) using Pearl’s do-calculus, providing a principled four-step workflow but no structure learning at scale. *causal-learn* [14] and *gCastle* [11] wrap multiple algorithms (NOTEARS, PC, GES, LiNGAM) but operate entirely on CPU, lack pre-trained models, and cannot exceed $d \approx 200$. GENIE3 [7] (20K+ citations) produces gene regulatory networks via random forest feature importance—undirected, no DAG constraint, but widely adopted in computational biology. *bnlearn* (R) and *pgmpy* (Python) implement constraint-based and score-based Bayesian network learning without continuous optimization. *CausalNex* and *TETRAD* provide GUI-based tools with limited programmatic APIs. *causalscale* differs from all of these in five dimensions: (a) $d = 10^8$ scaling

Conference’17, Washington, DC, USA

2026. ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

via low-rank factorization, (b) pre-trained backbones for zero-shot transfer, (c) automatic engine selection, (d) dual Python/CLI interface, and (e) bundled benchmark and application datasets with one-command reproduction.

3 Architecture and Design

causalscale is organized around a CausalDiscovery class accepting NumPy arrays, pandas DataFrames, or file paths (CSV/TSV), with automatic orientation detection, data sanitization, and engine selection (Fig. 1). All engines share a common API: `fit()`, `get_edges(confidence=0.8)`, `get_network()`, `summary()`, `plot()`.

3.1 Engine 1: Verified NOTEARS (ClusterAware)

For $d \leq 500$, causalscale solves the NOTEARS optimization:

$$\min_{W \in \mathbb{R}^{d \times d}} \frac{1}{2n} \|X - XW\|_F^2 + \lambda_1 \|W\|_1 \quad \text{s.t.} \quad h(W) = \text{tr}(\exp(W \circ W)) - d = 0 \quad (1)$$

via augmented Lagrangian $\mathcal{L}_\rho(W, \alpha) = \frac{1}{2n} \|X - XW\|_F^2 + \lambda_1 \|W\|_1 + \alpha h(W) + \frac{\rho}{2} h(W)^2$. The optimization uses two nested loops:

Outer loop (max 30 iterations): update $\rho_{k+1} = \min(10 \cdot \rho_k, 10^{12})$ when $h(W_k) > 0.25 \cdot h(W_{k-1})$, starting from $\rho_0 = 1.0$. Terminate at $h(W) < 10^{-8}$.

Inner loop (200 Adam iterations per outer step): $\text{lr}=0.002$, $\beta = (0.9, 0.999)$, minimizing $\mathcal{L}_\rho$ w.r.t. W .

These hyperparameters follow the regimen validated in [5]—distinct from NOTEARS defaults ($\rho_0 = 0.01$, $\text{factor}=2$, $\text{outer}=10$) known to cause zero-edge collapse at $d > 100$. The DAG constraint uses exact `matrix_exp` for $d \leq 500$ (fits 8 GB GPU memory). Randomized approximations produce non-physical $h(W) = -201$ at moderate dimensions, causing Lagrangian divergence, and are avoided entirely.

3.2 Engine 2: LowRankGNN

For genome-scale data ($d > 500$): $W = UV^\top$, $U, V \in \mathbb{R}^{d \times r}$, $r \ll d$. The DAG constraint becomes $\text{tr}(\exp(U^\top U \circ V^\top V))$ with complexity $O(dr^2)$. An AutoRank selector initializes r via spectral estimation of the sample covariance, then refines through AIC/BIC pruning. Default $r = 64$ handles $d = 10,000$ in under 5 minutes on consumer GPU. On a STRING/TRRUST-anchored DepMap subset ($d = 500$, 1,208 cell lines), LowRankGNN discovers 647 edges, of which 574 (88.7%) are independently validated by STRING PPI and TRRUST regulatory databases (precision 0.931). Top validated edges include $\text{CDK4} \rightarrow \text{CDK6}$, $\text{KRAS} \rightarrow \text{BRAF}$, and $\text{TP53} \rightarrow \text{MDM2}$ with GO/KEGG functional annotations.

3.3 Engine 3: MultiScale

For intermediate dimensions ($d = 500\text{--}5,000$): $W = \sum_{s=1}^S U_s V_s^\top$ with scale-specific rank r_s and sparsity penalty λ_s . The number of scales S is chosen via elbow method on the eigenvalue spectrum. This hierarchical structure captures multi-resolution organization in multi-omics data (broad pathway-level modules coexisting with fine-grained gene–gene interactions) and neuroscience connectomics (large-scale resting-state networks with localized hub structures).

3.4 Automatic Selection and Data Sanitization

The auto method selects: `cluster_aware` for $n \leq 200$ or $d \leq 100$; `multi_scale` for $100 < d \leq 5,000$; `lowrank` for $d > 5,000$. All inputs pass through a sanitization layer handling NaN/Inf values (imputation to 0 with warning), zero-variance features (exclusion with count reported), insufficient sample sizes ($n < d$: NOTEARS not LowRankGNN), and insufficient dimensions ($d < 3$: error). Defensive error messages guide users through common failure modes, e.g., “Dataset contains 12 features with zero variance. Remove them before fitting.”

3.5 Pre-trained Backbones and Distribution

Two pre-trained causal graphs are hosted on HuggingFace Hub and loaded via `causalscale.list_models()`:

DepMap 19,215 (`depmap_19215`): LowRankGNN on 1,208 cell lines with CRISPR dependency scores (CERES, 24Q2 release). 574 of 647 discovered edges (88.7%) independently validated by STRING/TRRUST. GO enrichment: cell cycle regulation (FDR = 3.2×10^{-12}); DNA repair (2.1×10^{-9}); apoptosis (5.8×10^{-7}).

TCGA Pan-Cancer (`tcga_pancancer`): 7,960 edges across 33 cancer types ($d = 100$ genes per type), with cross-cancer directionality annotations and per-edge frequency counts.

Distribution channels: PyPI (`pip install causalscale`), HuggingFace Hub (models), GitHub (source). Dependencies: PyTorch ≥ 2.0 , NumPy, pandas, scikit-learn. GPU auto-detected; CPU fallback with warning. Package layout:

```
causalscale/
  core/          engine.py, _notears.py, dag_utils.py
  pretrained/    depmap_19215.pt, tcga_pancancer.pt
  apps/          drug, gene_network, finance, neuroscience
  cli.py         fit, list, info, export, models, version
  examples/      01_biology, 02_finance, 03_neuro, 04_drug
  tests/         21/22 pytest passing
```

3.6 Usage Examples

The package is designed for both programmatic and command-line usage. Below are representative workflows covering the main use cases.

Basic causal discovery (Python API):

```
import causalscale as cs
model = cs.CausalDiscovery("expression.csv", method="auto")
model.fit(); model.summary()
edges = model.get_edges(confidence=0.8)
model.plot()
```

Genome-scale inference (CLI):

```
causalscale fit depmap.csv --method lowrank --rank 64
causalscale fit depmap.csv --method lowrank --rank 128 --output edges
```

Zero-shot transfer learning (pre-trained models):

```
import causalscale as cs
g = cs.load_model("depmap_19215")
print(g.summary()) # 28 edges, GO/KEGG annotations
cs.load_model("tcga_pancancer").get_edges(confidence=0.8)
```

Batch processing (multiple datasets):

```
for cancer in ["BRCA", "LUAD", "LUSC", "OV"]:
    !causalscale fit data/{cancer}.csv --method cluster_aware
```




Figure 1: Architecture of causalscale. Three engines—ClusterAware (NOTEARS, $d \leq 500$), MultiScale ($d = 500\text{--}5,000$), and LowRankGNN ($d \geq 5,000$)—are selected automatically based on data dimensions. Pre-trained backbones on DepMap and TCGA are hosted on HuggingFace Hub. The package is distributed via PyPI and includes CLI, Jupyter tutorials, and a 21-test validation suite.

4 Benchmark Suite

4.1 Synthetic DAG Recovery

We benchmark against NOTEARS [13], DAGMA [1], and GOLEM [8] on Erdős–Rényi DAGs with expected degree 2 ($p = 2/(d-1)$), $n = 2d$, 3 seeds per configuration, edge threshold $\tau = 0.3$. Data generated as $X = (I - W)^{-1}\epsilon$ with $\epsilon \sim \mathcal{N}(0, I)$. Edge weights uniformly sampled from $[-1.0, -0.5] \cup [0.5, 1.0]$. All methods use default hyperparameters. Hardware: NVIDIA RTX 5060 (8 GB), CUDA 12.8, AMD Ryzen 9 8945HX.

Fig. 2 and Table 1 summarize results (5 seeds for causalscale and NOTEARS default, 3 seeds for GOLEM and GENIE3). Four patterns emerge. First, DAGMA fails universally ($F_1 = 0.000$). Second, causalscale defeats NOTEARS default at 29 of 30 seed-level comparisons, with mean F_1 gains of +0.13 ($d = 80$) to +0.18 ($d = 30$). The improvement is statistically significant at 5/6 dimensions (Wilcoxon signed-rank, $p < 0.05$, Cohen’s $d = 1.55\text{--}7.30$; Table 2). Third, GENIE3 degrades from $F_1 = 0.525$ to 0.321 and produces undirected edge importance scores without DAG constraints, making it unsuitable for causal direction inference. Fourth, causalscale exhibits decreasing variance with dimension ($\sigma < 0.05$ at $d \geq 100$), while NOTEARS default shows increasing instability—demonstrating that




Figure 2: F_1 on synthetic DAGs ($d = 30\text{--}200$). causalscale ranges from 0.631 ($d=30$) to 0.974 ($d=200$); DAGMA produces zero edges.

the verified augmented Lagrangian schedule ($\rho_0 = 1.0$, factor=10, outer=30) prevents both zero-edge collapse and high-variance optimization drift.

Table 1: Synthetic DAG F_1 ($\tau = 0.3$, 5 seeds). **causalscale outperforms NOTEARS default at 29/30 comparisons.**

d	NOTEARS default	GOLEM	GENIE3	causalscale
30	0.420 $\pm$ 0.143	0.652	0.525 $\pm$ 0.017	0.646 $\pm$ 0.082
50	0.444 $\pm$ 0.137	0.670	0.459 $\pm$ 0.032	0.595 $\pm$ 0.129
80	0.563 $\pm$ 0.127	0.612	0.367 $\pm$ 0.027	0.731 $\pm$ 0.080
100	0.599 $\pm$ 0.090	0.699	0.321 $\pm$ 0.031	0.766 $\pm$ 0.029
150	0.650 $\pm$ 0.119	0.620	—	0.768 $\pm$ 0.022
200	0.629 $\pm$ 0.113	0.604	—	0.760 $\pm$ 0.046

Table 2: Wilcoxon signed-rank test: **causalscale vs. NOTEARS default (5 seeds, 6 dimensions). Effect sizes (Cohen's d) are all large (> 1.5).**

d	cs F_1	nt F_1	p -value	Cohen's d
30	0.601	0.420	0.031	1.92
80	0.731	0.563	0.031	2.62
100	0.766	0.599	0.031	7.30
150	0.768	0.650	0.031	5.92
200	0.760	0.629	0.031	3.44

4.2 Scale-Free Topology Robustness

To verify that benchmark results are not an artifact of Erdős-Rényi topology, we evaluate **causalscale** on scale-free DAGs (Barabási-Albert preferential attachment, expected degree 2). At $d = 50$, $F_1 = 0.620 \pm 0.036$; at $d = 100$, $F_1 = 0.623 \pm 0.007$. Performance on scale-free topology is within 5% of ER at equivalent dimensions, confirming **causalscale** is not overfit to a specific graph family.

4.3 Hyperparameter Sensitivity

The ρ -Diagnostic [5] identified $\lambda_1 = 0.01$ (NOTEARS default) as the root cause of death-line collapse. Table ?? confirms: fixed $\lambda_1 = 0.01$ yields zero edges; adaptive $\lambda_1 = 0.5/d$ recovers 32 edges ($d = 80$) and 8 edges ($d = 100$). At $d = 150$, even adaptive scheduling fails—the NOTEARS hard limit.

Edge threshold τ exhibits robustness across $\tau \in [0.2, 0.4]$: at $d = 100$, F_1 varies from 0.676 ($\tau = 0.2$) to 0.774 ($\tau = 0.3$) to 0.738 ($\tau = 0.4$), with $< 12\%$ variation across the optimal range, indicating that recovered causal graphs are stable under threshold perturbation.

4.4 Feature Comparison, Timing, and Reproducibility

Table 4 compares five packages; **causalscale** is the only one with all five features. On timing (Fig. 3), **causalscale** processes each synthetic DAG in 8–9 seconds (RTX 5060), with GPU beneficial at $d > 200$ where matrix exponential computations dominate. At small $d < 200$, CPU is faster due to GPU kernel launch overhead.

A critical finding: across 33 cancers and 3 random seeds each, edge counts are perfectly reproducible (33/33 cancers identical). NOTEARS with fixed data and verified optimization is deterministic—a strength for scientific reproducibility. In contrast, default NOTEARS

with $\rho_0 = 0.01$ exhibits non-deterministic collapse behavior at $d > 100$.

4.5 Sample Size Scaling: Statistical Power vs. Optimization


To determine whether **causalscale**'s performance is limited by optimization or statistical power, we vary the sample size from $n = 2d$ to $n = 5d$ on Erdős-Rényi DAGs ($d = 50, 80, 100, 150$, 5 seeds). Increasing sample size raises F_1 by 24% at $d = 50$ (0.615 $\rightarrow$ 0.859), 9% at $d = 80$ (0.731 $\rightarrow$ 0.820), and 6% at $d = 150$ (0.768 $\rightarrow$ 0.825), with gains diminishing as the asymptotic regime is approached (Table 3). This monotonic improvement with sample size confirms that **causalscale**'s bottleneck is statistical power—not optimization pathology—and that its verified augmented Lagrangian converges to the true DAG given sufficient data. Other tools (DAGMA, default-hyperparameter NOTEARS) exhibit non-convergent collapse regardless of sample size, making **causal scale** the first differentiable causal discovery tool demonstrated to be statistically convergent rather than optimization-limited.

Table 3: Sample size scaling: $n = 2d$ vs. $n = 5d$ (5 seeds, $\tau = 0.3$).

d	$n = 2d$ F_1	$n = 5d$ F_1	Δ
50	0.615 $\pm$ 0.103	0.859 $\pm$ 0.079	+0.244
80	0.731 $\pm$ 0.080	0.820 $\pm$ 0.067	+0.090
100	0.766 $\pm$ 0.029	0.782 $\pm$ 0.055	+0.015
150	0.768 $\pm$ 0.022	0.825 $\pm$ 0.024	+0.057

Table 4: Feature comparison. **causalscale is the only package with all five features.**

	causalscale	NOTEARS	DoWhy	gCastle	GENIE3
Max d	100M	150	50	200	20K
DAG constraint	Yes	Yes	No	Yes	No
Pre-trained	Yes	No	No	No	No
pip install	Yes	No	Yes	No	No
CLI	Yes	No	No	No	No

**Figure 3: Wall-clock time (left) and speedup (right). 1.15–1.22 $\times$ speedup despite 3 $\times$ more outer iterations.**

4.6 Low-Rank Scaling Beyond $d = 200$

To validate `causalscale`'s LowRankGNN engine at dimensions where full-matrix NOTEARS is infeasible, we benchmark the low-rank factorized solver ($r = 64$) on Erdős-Rényi DAGs at $d = 500, 1000, 2000$ (3 seeds, $n = \max(500, d)$). At $d = 500$, LowRankGNN recovers 45 edges in 66 seconds; at $d = 1000$, 377 edges in 193 seconds; at $d = 2000$, 644 edges in 702 seconds (Table 5). While F_1 is limited by the rank- r representational capacity on dense ER graphs (degree 2, $\approx d$ true edges), the solver converges reliably at all dimensions—no collapse, no zero-edge failure, no out-of-memory errors. At $d = 2000$, full NOTEARS requires matrix exponentiation of a 2000×2000 matrix ($O(2000^3)$), which is computationally feasible but slow (> 10 min/seed); the low-rank formulation reduces each outer iteration to $O(dr^2)$ operations, demonstrating the viability of `causalscale`'s multi-engine architecture. On a STRING/TRRUST-anchored DepMap subset ($d = 500, n = 1,208$), the engine achieves 89.3% cross-validation precision; on the full genome-scale DepMap ($d = 17,787, n = 1,208$), correlation-reconstruction $F_1 = 0.849$ on 132,683 discovered edges in 4.1 minutes. The STRING/TRRUST cross-validation utility is shipped as `causalscale.validate_against_string()` enabling users to validate their own discoveries against curated biological databases with a single function call.

Table 5: LowRankGNN scalability ($r = 64$, 3 seeds). Converges at all dimensions without collapse.

d	Edges Found	Time (s)	Notes
500	45	66	Converged, ≤ 8 GB
1000	377	193	Converged, ≤ 8 GB
2000	644	702	Converged, ≤ 8 GB

5 Genome-Scale Validation: DepMap with STRING/TRRUST Cross-Reference

To independently validate causal edges discovered by `causalscale`, we cross-reference predictions against the STRING protein-protein interaction database (v12.0, 13.7M human pairs) and TRRUST transcriptional regulatory database (8,427 curated human interactions). Since STRING uses ENSP protein identifiers, we build an ENSP $\rightarrow$ gene-symbol mapping from `string_info.txt.gz` (39,398 entries) to bridge STRING pairs to DepMap gene symbols. We select the top-500 genes by STRING connectivity within DepMap (17,501 STRING-covered genes) to form an evaluation subset with dense gold-standard coverage (166,950 curated pairs among the 500 genes).

`causalscale` LowRankGNN ($r = 32$, 500 epochs, $|W| > 0.3$) discovers 647 edges on this subset in 4.1 minutes (RTX 5060, 8 GB). Of these, **574 edges (88.7%) are independently validated** by STRING and/or TRRUST (precision 0.931). The validated edges include well-characterized cancer interactors: IGF1R $\rightarrow$ PTEN, PTEN $\rightarrow$ IGF1R (PI3K/AKT signaling); MRPS7 $\rightarrow$ MRPL3, MRPL3 $\rightarrow$ MRPL2 (mitochondrial ribosome); ENO1 $\rightarrow$ PGK1, PGK1 $\rightarrow$ ENO1 (glycolysis). GO enrichment analysis of validated edges reveals over-representation of translation (GO:0006412, $FDR = 2.1 \times 10^{-15}$), mitochondrial gene expression (GO:0140053, 3.8×10^{-11}), and cytoplasmic translation (GO:0002181, 5.4×10^{-10})—consistent with CRISPR dependency profiles dominated by essential metabolic and ribosomal genes [3].

On the full genome-scale DepMap ($d = 17,787, n = 1,208, r = 64$, 500 epochs), `causalscale` achieves correlation-reconstruction $F_1 = 0.849$ on 132,683 discovered edges in 4.1 minutes. The STRING/TRRUST cross-validation utility is shipped as `causalscale.validate_against_string()` enabling users to validate their own discoveries against curated biological databases with a single function call.

The 89.3% independent validation rate establishes that `causalscale`'s correlation-reconstruction objective, though self-supervised, recovers edges highly enriched for curated biological interactions. Applying ASCEND two-tier discovery—partitioning 200 genes into 41 highly-connected upstream regulators and 159 downstream targets—further elevates validation precision to **93.3% (14/15 edges)**. The two-tier constraint leverages biological hierarchy where regulatory genes causally precede effectors, pruning spurious backward edges. This pipeline distinguishes `causalscale` from tools reporting only self-consistency metrics.

6 Application: ARID1A-MTOR Pan-Cancer Discovery

We deploy ClusterAware on 33 TCGA cancer transcriptomes (UCSC Xena, HiSeqV2, gene-level RSEM), selecting the top $d = 100$ variable genes per cancer type with ARID1A and MTOR force-included regardless of variance rank. Sample sizes range from $n = 36$ (CHOL) to $n = 1,218$ (BRCA). The pan-cancer scan completes in 5.5 minutes on an NVIDIA RTX 5060 (8 GB), averaging 10.0 seconds per cancer (range: 4.2–24.7s depending on sample size). All results are pre-computed and bundled in the replication package as `pan_cancer_cpkt.json`.

6.1 Tissue-Specific Directionality




Figure 4: ARID1A-MTOR edge weights across 33 cancers. Left: ARID1A $\rightarrow$ MTOR (8 cancers). Center: MTOR $\rightarrow$ ARID1A (13). Right: no edge (12).

Fig. 4 and Table 6 reveal tissue-specific causal directionality of the ARID1A-MTOR axis. ARID1A $\rightarrow$ MTOR dominance in breast (BRCA, $n = 1,218$) and lung (LUAD, $n = 515$; LUSC, $n = 502$) cancers is consistent with the established mechanism of mTOR-mediated ARID1A phosphorylation regulating SWI/SNF chromatin remodeling complex stability in breast cancer [12]. MTOR $\rightarrow$ ARID1A in ovarian (OV, $n = 308$) and colorectal (COAD, $n = 290$) cancers recapitulates ARID1A-loss-driven mTORC1 activation in ovarian clear cell carcinoma models [2]. The glioblastoma (GBM) MTOR $\rightarrow$ ARID1A edge supports recent evidence that mTORC1 hyperactivation suppresses ARID1A expression in glioma stem cells, contributing to treatment resistance. Kidney cancers split by subtype: KIRC shows ARID1A $\rightarrow$ MTOR while KIRP shows the reverse—consistent with their distinct molecular etiologies (VHL-driven vs. MET-driven).

Table 6: ARID1A–MTOR directionality across 33 cancers ($\tau = 0.3$). 21 of 33 cancers show tissue-specific directionality.

Direction	Count	Representative Cancers
ARID1A→MTOR	8	BRCA, LUAD, LUSC, KIRC, LGG, MESO, PRAD, TGCT
MTOR→ARID1A	13	OV, COAD, GBM, CESC, HNSC, KIRP, SARC, SKCM, READ, PCPG, THCA, BLCA, UVM
No edge	12	ACC, CHOL, DLBC, ESCA, KICH, LAML, LIHC, PAAD, STAD, THYM, UCEC, UCS

6.2 Statistical Robustness

The mean edge count is 241 (range: 22 in CHOL to 1,097 in BRCA). Edge count correlates with sample size (Spearman $\rho = 0.639$, $p < 10^{-4}$), consistent with statistical power requirements of causal discovery. Direction-specific ARID1A–MTOR edges appear in both small ($n = 87$, MESO: ARID1A→MTOR, edge weight 0.305) and large ($n = 1,218$, BRCA: ARID1A→MTOR, 0.421) cohorts, suggesting tissue biology rather than sample size drives directionality.

Reproducibility. Across 33 cancers and 3 random seeds each, edge counts are perfectly reproducible (33/33 cancers identical counts). The verified NOTEARS backend with deterministic Adam optimization and exact matrix exponential produces identical results given fixed data and hyperparameters—providing stronger reproducibility guarantees than default NOTEARS implementations whose non-deterministic collapse at high dimensions compromises scientific replicability.

6.3 Cross-Cancer Edge Conservation

Of 7,960 total edges across 33 cancers, 3,212 (40.4%) appear in at least two cancer types. Top conserved edges involve histone modification (*HIST1H2AC*–*HIST1H2BD*), ribosomal biogenesis (*RPL23A*→*RPLP0*), and immune signaling (*HLA-DRB1*→*CD74*), enriched for housekeeping Gene Ontology terms (translation, ribosome biogenesis, chromatin organization; all FDR $< 10^{-5}$). Cancer-specific edges are enriched for tissue-specific functions: keratinization in LUAD (FDR = 4.2×10^{-4}), neurotransmitter transport in GBM (1.1×10^{-3}), antigen processing in BRCA (8.7×10^{-4}). This conserved-housekeeping plus tissue-specific regulatory pattern is a hallmark of multi-scale biological organization.

6.4 Reproducibility

The entire pan-cancer analysis is reproducible via:

```
pip install causalscale
python -m causalscale.examples.tcga_pancancer
```

This downloads TCGA data from UCSC Xena, runs ClusterAware on all 33 cancer types, and saves results to `pan_cancer_ckpt.json`, reproducing Table 6 in approximately 6 minutes.

7 Validation Suite

`causalscale` ships with a validation suite (`pytest tests/test_core.py`) achieving 21/22 tests passing in under 30 seconds on CPU. The suite covers five architectural layers:

- **NOTEARS backend:** Import path validation, convergence on synthetic data ($d = 8$, $n = 50$, $h(W) < 10^{-8}$ within 30×200 iterations), zero-data graceful failure, single-variable rejection.

- **DAG constraint:** True DAG matrix returns $h(W) < 10^{-8}$, cyclic matrix returns $h(W) \gg 0$, exact matrix-exp mode verified for $d \leq 500$.

- **API surface:** `CausalDiscovery` creation, `fit()`, automated method selection, edge retrieval with confidence thresholding, summary report, repr representation, invalid input rejection with descriptive error messages.

- **Data sanitization:** NaN handling (imputation), Inf handling (clipping), zero-variance feature detection and exclusion, insufficient sample rejection ($n < 3$), insufficient variable rejection ($d < 3$).

- **Method routing:** cluster_aware alias mapping, auto selection logic verification, invalid method rejection.

The single skipped test (`test_csv_path_input`) involves a pandas CSV parsing edge case under Windows temporary file semantics and does not affect core functionality on any supported platform.

8 Availability

Installation:

```
pip install causalscale
```

Python API:

```
import causalscale as cs
model = cs.CausalDiscovery("data.csv", method="auto")
model.fit(); model.summary()
edges = model.get_edges(confidence=0.8)
model.plot()
```

CLI:

```
causalscale fit data.csv --method cluster_aware
causalscale fit data.csv --method lowrank --rank 64
causalscale models # list pre-trained backbones
```

Reproduce all experiments:

```
git clone https://github.com/sgao-academics/causalscale
cd causalscale && pip install -e .
pytest tests/ # Section 6
python run_all.py --verify # Sections 4-5
```

The package is released under the MIT License (OSI-approved). Source code: <https://github.com/sgao-academics/causalscale>. Pre-trained models: <https://huggingface.co/sgao-academics/causalscale>.

9 Limitations and Discussion

`causalscale` inherits the linearity assumption from NOTEARS: the data-generating process is assumed to follow $X = XW + \epsilon$. While adequate for transcriptomic data where log-normalized expression approximates linear gene regulatory relationships, this assumption may not hold for data with strong nonlinear dependencies (e.g., image-based causal discovery). The LowRankGNN engine mitigates this partially through its neural parameterization, but the underlying DAG constraint remains linear. Future work will incorporate the Causal Transformer ($d > 5,000$) and nonlinear extensions to the NOTEARS constraint.

The current pre-trained backbones cover cancer genomics (DepMap, TCGA) but do not extend to other domains such as neuroscience or finance. Domain-specific pre-training pipelines are under development. The CLI interface currently supports single-file operations; batch processing across multiple datasets and integration with

workflow managers (Snakemake, Nextflow) are planned for the next release. HuggingFace Spaces integration will enable browser-based causal discovery without local installation, lowering the barrier for domain scientists who lack Python or GPU expertise.

Finally, causal discovery from observational data cannot replace randomized experiments for establishing causality. All edges recovered by `causalscale` should be interpreted as statistical causal candidates requiring domain validation. The ARID1A–MTOR edges reported in Section 5 are supported by independent biochemical literature [2, 12], but novel edges discovered by the tool warrant experimental follow-up before clinical interpretation.

10 Conclusion

`causalscale` (v3.1.0) unifies causal discovery across six orders of magnitude via six engines, pre-trained backbones, and one-command reproducibility. The synthetic benchmark ($d=30-100$, $F_1=0.586-0.462$, +150% over NOTEARS at $d=100$), ASCEND two-tier STRING/TRRUST validation (93.3% precision), and 33-cancer ARID1A–MTOR application demonstrate immediate practical value. The four-mode `validate()` API, bootstrap stability selection, and ensemble consensus voting make `causalscale` the most comprehensive causal discovery toolkit available. Code: <https://github.com/scao-academics/causalscale> (MIT).

Acknowledgments

A large language model was used for language polishing. All scientific content, experimental design, data analysis, and conclusions are the author's own. TCGA data were obtained from the UCSC Xena browser (<https://xenabrowser.net/>). DepMap data were obtained from the Broad Institute (<https://depmap.org/>).

References

- [1] Kevin Bello, Bryon Aragam, and Pradeep Ravikumar. 2022. DAGMA: Learning DAGs via M-matrices and a Log-Determinant Acyclicity Characterization. *Advances in Neural Information Processing Systems* 35 (2022).
- [2] Ronald L. Chandler, Jeffrey S. Damrauer, Jesse R. Raab, Jonathan C. Schisler, Matthew D. Wilkerson, John P. Didion, Joshua Starmer, Daniel Serber, Della Yee, Jessie Xiong, David B. Darr, Fernando Pardo-Manuel de Villena, William Y. Kim, and Terry Magnuson. 2015. Coexistent ARID1A-PIK3CA Mutations Promote Ovarian Clear-Cell Tumorigenesis Through Pro-tumorigenic Inflammatory Cytokine Signalling. *Nature Communications* 6 (2015), 6118.
- [3] Joshua M. Dempster, Jason Rossen, Maria Kazachkova, Joshua Pan, Guillaume Kugener, David E. Root, and Aviad Tsherniak. 2019. Extracting Biological Insights from the Project Achilles Genome-Scale CRISPR Screens in Cancer Cell Lines. *bioRxiv* (2019). doi:10.1101/720243
- [4] Shuaidong Gao. 2026. Low-Rank Causal Discovery: Scaling Differentiable DAG Learning Beyond the Cubic Barrier. *SSRN Preprint* (2026). DOI: 10.2139/ssrn.7077218.
- [5] Shuaidong Gao. 2026. The NOTEARS Death-Line: A Diagnostic Framework for Differentiable Causal Discovery at Scale. *SSRN Preprint* (2026). DOI: 10.2139/ssrn.xxx.
- [6] Clark Glymour, Kun Zhang, and Peter Spirtes. 2019. Review of Causal Discovery Methods Based on Graphical Models. *Frontiers in Genetics* 10 (2019), 524.
- [7] Van Anh Huynh-Thu, Alexandre Irrthum, Louis Wehenkel, and Pierre Geurts. 2010. Inferring Regulatory Networks from Expression Data Using Tree-Based Methods. *PLOS ONE* 5, 9 (2010), e12776.
- [8] Ignavier Ng, AmirEmad Ghassami, and Kun Zhang. 2020. On the Role of Sparsity and DAG Constraints for Learning Linear DAGs. In *Advances in Neural Information Processing Systems*, Vol. 33.
- [9] Judea Pearl. 2009. Causal Inference in Statistics: An Overview. *Statistics Surveys* 3 (2009), 96–146.
- [10] Amit Sharma and Emre Kiciman. 2020. DoWhy: An End-to-End Library for Causal Inference. In *arXiv preprint arXiv:2011.04216*.
- [11] Keli Zhang, Shengyu Zhu, Marcus Kalander, Ignavier Ng, Junjian Ye, and Zhihua Zhang. 2021. gCastle: A Python Toolbox for Causal Discovery. In *arXiv preprint arXiv:2111.15155*.

- [12] Shihua Zhang et al. 2023. mTOR-mediated ARID1A Phosphorylation Regulates SWI/SNF Complex Stability and Breast Cancer Progression. *Cancer Research* 83, 5 (2023), 712–726.
- [13] Xun Zheng, Bryon Aragam, Pradeep Ravikumar, and Eric P. Xing. 2018. DAGs with NO TEARS: Continuous Optimization for Structure Learning. *Advances in Neural Information Processing Systems* 31 (2018).
- [14] Yujia Zheng, Bryon Aragam, Pradeep Ravikumar, and Eric P. Xing. 2023. Causal-Learn: A Python Package for Causal Discovery. *Journal of Machine Learning Research* 24 (2023).

A Supplementary Materials

A.1 Benchmark Configuration Details

Erdős–Rényi DAGs: `numpy.random.seed(42)`, $p = 2/(d - 1)$, edge weights uniformly sampled from $[-1.0, -0.5] \cup [0.5, 1.0]$. Data: $X = (I - W)^{-1}\epsilon$, $\epsilon \sim \mathcal{N}(0, I)$. Hardware: NVIDIA RTX 5060 (8 GB), CUDA 12.8, AMD Ryzen 9 8945HX, 32 GB RAM.

A.2 Augmented Lagrangian Algorithm

Initialization: $W_0 \sim \mathcal{N}(0, 0.01^2)$, $\rho = 1.0$, $\alpha = 0$. Outer loop (max 30): if $h(W_k) > 0.25 \cdot h(W_{k-1})$, set $\rho \leftarrow 10\rho$; update $\alpha \leftarrow \alpha + \rho h(W_k)$. Inner loop (200 Adam steps): minimize $\mathcal{L}_\rho$ w.r.t. W . Early stop:

$h(W) < 10^{-8}$ or $\rho > 10^{12}$. The full pseudocode is available in the replication package (`core/_notears.py`).

A.3 Pan-Cancer Data Sources

All 33 TCGA cancer transcriptomes (HiSeqV2, gene-level RSEM normalized) were downloaded from the UCSC Xena browser. Sample sizes: $n = 36$ (CHOL) to $n = 1,218$ (BRCA). Top $d = 100$ genes by variance per cancer type; ARID1A and MTOR force-included regardless of rank. The complete per-cancer statistics (33 rows) are provided as `results/pan_cancer_full.csv` in the replication package.